

# Introduction to R

Day 1 - Introduction to Data Analysis with R

Selina Baldauf

Freie Universität Berlin - Theoretical Ecology

October 17, 2023

# R as a calculator

## Arithmetic operators

Addition	+
Subtraction	-
Multiplication	*
Division	/
Modulo	%%
Power	^

```
# Addition
2 + 2
# Subtraction
5.432 - 34234
# Multiplication
33 * 42
# Division
3 / 42
# Modulo (Remainder)
2 %% 2
# Power
2^2
# Combine operations
((2 + 2) * 5)^ (10 %% 10)
```

# R as a calculator

## Relational operators

Equal to	<code>==</code>	<pre>2 == 2 #&gt; [1] TRUE</pre>
Not equal to	<code>!=</code>	<pre>2 != 2 #&gt; [1] FALSE</pre>
Less than	<code>&lt;</code>	<pre>33 &lt;= 32 #&gt; [1] FALSE</pre>
Greater than	<code>&gt;</code>	<pre>20 &lt; 20 #&gt; [1] FALSE</pre>
Less or equal than	<code>&lt;=</code>	
Greater or equal than	<code>&gt;=</code>	

# R as a calculator

## Logical operators

Not

---

---

!

```
!TRUE  
#> [1] FALSE  
!(3 < 1)  
#> [1] TRUE
```

# R as a calculator

## Logical operators

Not	!
And	&

```
(3 < 1) & (3 == 3) # FALSE & TRUE = FALSE
#> [1] FALSE
(1 < 3) & (3 == 3) # TRUE & TRUE = TRUE
#> [1] TRUE
(3 < 1) & (3 != 3) # FALSE & FALSE = FALSE
#> [1] FALSE
```

# R as a calculator

## Logical operators


Not	!
And	&
Or	

```
(3 < 1) | (3 == 3) # FALSE | TRUE = TRUE
#> [1] TRUE
(1 < 3) | (3 == 3) # TRUE | TRUE = TRUE
#> [1] TRUE
(3 < 1) | (3 != 3) # FALSE | FALSE = FALSE
#> [1] FALSE
```

# Basic R Syntax

- Whitespace does not matter

```
# this  
data<-read_csv("data/my-data.csv")  
  
# is the same as this  
  
data <-  
  read_csv(    "data/my-data.csv"    )
```

- There are good practice rules however -> More on that later
- RStudio will (often) tell you if something is incorrect
  - Find  on the side of your script

# Comments in R

```
# Reading and cleaning the data -----  
  
data <- read_csv("data/my-data.csv")  
# clean all column headers  
# (found on https://stackoverflow.com/questions/68177507/)  
data <- janitor::clean_names(data)  
  
# Analysis -----
```

- Everything that follows a # is a comment
- Comments are not evaluated
- Notes that make code more readable or add information
- Comments can be used for
  - Explanation of code (if necessary)
  - Include links, names of authors, ...
  - Mark different sections of your code (💡 try `Ctrl/Cmd + Shift + R`)



# Objects and data types in

# Variables

- Store values under meaningful names **to reuse** them
- A variable has a **name** and **value** and is created using the **assignment operator**

`radius <- 5`

- Variables are available in the global environment
- R is case sensitive: `radius` != `Radius`
- Variables can hold any R objects, e.g. numbers, tables with data, ...
- Choose meaningful variable names
  - Make your code easier to read

# Variables

```
# create a variable
radius <- 5
# use it in a calculation and save the result
# pi is a built-in variable that comes with R
circumference <- 2 * pi * radius
# change value of variable radius
radius <- radius + 1
```

```
# just use the name to print the value to the console
radius
```

# Atomic data types

There are 6 so-called **atomic data types** in R. The 4 most important are:

**Numeric:** There are two numeric data types:

- **Double:** can be specified in decimal (`1.243` or `-0.2134`), scientific notation (`2.32e4`) or hexadecimal (`0xd3f1`)
- **Integer:** numbers that are not represented by fraction. Must be followed by an `L` (`1L`, `2038459L`, `-5L`)

**Logical:** only two possible values `TRUE` and `FALSE` (abbreviation: `T` or `F` - but better use non-abbreviated form)

**Character:** also called string. Sequence of characters surrounded by quotes (`"hello"` , `"sample_1"`)

# Vectors

Vectors are data structures that are built on top of atomic data types.

Imagine a vector as a **collection of values** that are all of the same data type.




Image from [Advanced R book](#)

# Creating vectors

Use the function `c()` to *combine* values into a vector

```
lgl_var <- c(TRUE, TRUE, FALSE)
dbl_var <- c(2.5, 3.4, 4.3)
int_var <- c(1L, 45L, 234L)
chr_var <- c("These are", "just", "some strings")
```

The `:` operator creates a sequence between two numbers with an increment of (-)1

```
1:10 # instead of c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
#> [1] 1 2 3 4 5 6 7 8 9 10
```

There are many more options to create vectors

- `seq()` to create a sequence of numbers
- `rep()` to repeat values
- ...

# Creating vectors: `c()`

Be aware of implicit **type conversion** when combining vectors of different types

```
# integer + logical -> integer (same with double + logical)
c(int_var, lgl_var)
#> [1] 1 45 234 1 1 0

# integer + character -> character (same with double + character)
c(int_var, chr_var)
#> [1] "1" "45" "234" "These are" "just"
#> [6] "some strings"

# logical + character -> character
c(lgl_var, chr_var)
#> [1] "TRUE" "TRUE" "FALSE" "These are" "just"
#> [6] "some strings"
```

# Working with vectors



# Working with vectors

Let's create some vectors to work with.

```
# list of 10 biggest cities in Europe
cities <- c("Istanbul", "Moscow", "London", "Saint Petersburg", "Berlin",
            "Madrid", "Kyiv", "Rome", "Bucharest", "Paris")

population <- c(15.1e6, 12.5e6, 9e6, 5.4e6, 3.8e6, 3.2e6, 3e6, 2.8e6, 2.2e6, 2.1e6)

area_km2 <- c(2576, 2561, 1572, 1439, 891, 604, 839, 1285, 228, 105 )
```

We can check the length of a vector using the `length()` function:

```
length(cities)
#> [1] 10
```

# Working with vectors

Divide population and area vector to calculate population density in each city:

```
population / area_km2
#> [1] 5861.801 4880.906 5725.191 3752.606 4264.871 5298.013 3575.685
#> [8] 2178.988 9649.123 20000.000
```

The operation is performed **separately for each element of the two vectors** and the result is a vector.

Same, if a **vector is divided by vector of length 1** (i.e. a single number). Result is always a vector.

```
mean_population <- mean(population) # calculate the mean of population vector
mean_population
#> [1] 5910000
population / mean_population # divide population vector by the mean
#> [1] 2.5549915 2.1150592 1.5228426 0.9137056 0.6429780 0.5414552 0.5076142
#> [8] 0.4737733 0.3722504 0.3553299
```

# Working with vectors

We can also work with relational and logical operators

```
population > mean_population  
#> [1] TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE
```

The result is a vector containing **TRUE** and **FALSE**, depending on whether the city's population is larger than the mean population or not.

Logical and relational operators can be combined

```
# population larger than mean population OR population larger than 3 million  
population > mean_population | population > 3e6  
#> [1] TRUE TRUE TRUE TRUE TRUE TRUE FALSE FALSE FALSE FALSE
```

# Working with vectors

Check whether elements occur in a vector:

```
cities == "Istanbul"
#> [1] TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

The `%in%` operator checks whether *multiple* elements occur in a vector.

```
# for each element of cities, checks whether that element is contained in to_check
to_check <- c("Istanbul", "Berlin", "Madrid")
cities %in% to_check # same as cities %in% c("Istanbul", "Berlin", "Madrid")
#> [1] TRUE FALSE FALSE FALSE TRUE TRUE FALSE FALSE FALSE FALSE
```

`%in%` always returns a vector of the same length as the vector on the left side

```
# for each element of to_check, check whether that element is contained in cities
to_check %in% cities
#> [1] TRUE TRUE TRUE
```

# Indexing vectors

You can use square brackets `[]` to access specific elements from a vector.

The basic structure is:

`vector [ vector of indexes to select ]`

```
cities[5]  
#> [1] "Berlin"
```

```
# the three most populated cities  
cities[1:3] # same as cities[c(1,2,3)]  
#> [1] "Istanbul" "Moscow"   "London"
```

```
# the last entry of the cities vector  
cities[length(cities)] # same as cities[10]  
#> [1] "Paris"
```

# Indexing vectors

Change the values of a vector at specified indexes using the assignment operator `<-`

Imagine for example, that the population of

- Istanbul (index 1) increased to 20 Million
- Rome (index 8) changed but is unknown
- Paris (index 10) decreased by 200,000

```
# Update Istanbul (1) and Rome(8)
population[c(1, 8)] <- c(20e6, NA) # NA means missing value
# Update Paris (10)
population[10] <- population[10] - 200000

# Look at the result
population
#>  [1] 20000000 12500000  9000000  5400000  3800000  3200000  3000000          NA
#>  [9]  2200000  1900000
```

# Indexing vectors

You can also index a vector using logical tests. The basic structure is:

**vector [ logical vector of same length ]**

```
mega_city <- population > mean_population  
mega_city  
#> [1] TRUE TRUE TRUE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

Which are the mega cities?

```
cities[mega_city] # or short: cities[population > mean_population]  
#> [1] "Istanbul" "Moscow" "London"
```

Return only the cities for which the comparison of their population against the mean population is **TRUE**

# Summary

Introduction to R



# Summary I

- Variables have a name and a value and are created using the assignment operator `<-`, e.g.

```
radius <- 5
```

- Vectors are a collection of values of the same data type:
  - character (`"hello"`)
  - numeric: integer (`23L`) and double (`2.23`)
  - logical (`TRUE` and `FALSE`)

# Summary II

## Create vectors

```
# combine objects into vector
c(1,2,3)

# create a sequence of values
seq(from = 3, to = 6, by = 0.5)
seq(from = 3, to = 6, length.out = 10)
2:10

# repeat values from a vector
rep(c(1,2), times = 2)
rep(c("a", "b"), each = 2)
```

# Summary III

## Indexing and subsetting vectors

```
# By index
v[3]
v[1:4]
v[c(1, 5, 7)]

# Logical indexing with 1 vector
v[v > 5]
v[v != "bird" | v == "rabbit"]
v[v %in% c(1, 2, 3)] # same as v[v == 1 | v == 2 | v == 3]

# Logical indexing with two vectors of same length
v[y == "bird"] # return the value in v for which index y == "bird"
v[y == max(y)] # return the value in v for which y is the maximum of y
```

# Summary IV

## Working with vectors

```
# length
length(v)
# rounding numbers
round(v, digits = 2)
# sum
sum(v)
# mean
mean(v)
# median
median(v)
# standard deviation
sd(v)
# find the min value
min(v)
# find the max value
```

# Now you

Task (35 min)

Working with vectors

Find the task description [here](#)

