

Data frames and tibbles: tables in R

Introduction to R - Day 1

Instructor: **Selina Baldauf**

Freie Universität Berlin - Theoretical Ecology

2021-08-01 (updated: 2022-08-26)

Data frames

The built-in data structure for tables in R is a `data frame`.

- Vectors in R can't represent a table with data that is connected via rows

Data frames are one of the **biggest and most important ideas** in R, and one of the things that make R different from other programming languages.

(Wickham, [Advanced R](#))

cities	population	area_km2
Istanbul	15100000	2576
Moscow	12500000	2561
London	9000000	1572
Saint Petersburg	5400000	1439
Berlin	3800000	891
Madrid	3200000	604
Kyiv	3000000	839
Rome	2800000	1285
Bucharest	2200000	228
Paris	2100000	105

Data frames

A data frame is a **named list of vectors** of the same length.

Basic properties of a data frame

- every **column is a vector**
- columns have a **header**
 - this is the **name** of the vector in the list
- within one column, all values are of the **same data type**
- every column has the same length

The diagram shows a table representing a data frame. The first row is the header row, containing the column names: 'cities', 'population', and 'area_km2'. The subsequent rows contain data for various cities. Annotations include a yellow arrow pointing to the 'cities' header with the label 'character', and two yellow arrows pointing to the 'population' and 'area_km2' headers with the label 'numeric'. A blue arrow points downwards along the left side of the table, indicating the rows.

cities	population	area_km2
Istanbul	15100000	2576
Moscow	12500000	2561
London	9000000	1572
Saint Petersburg	5400000	1439
Berlin	3800000	891
Madrid	3200000	604
Kyiv	3000000	839
Rome	2800000	1285
Bucharest	2200000	228
Paris	2100000	105

Data frames

Data frames are created with the function `data.frame()`:

```
cities <- c(
  "Istanbul", "Moscow", "London",
  "Saint Petersburg", "Berlin", "Madrid",
  "Kyiv", "Rome", "Bucharest", "Paris")

population <- c(
  15.1e6, 12.5e6, 9e6, 5.4e6, 3.8e6,
  3.2e6, 3e6, 2.8e6, 2.2e6, 2.1e6)

area_km2 <- c(2576, 2561, 1572, 1439,
  891, 604, 839, 1285, 228, 105)
```

```
data.frame(
  cities = cities,
  population = population,
  area_km2 = area_km2
)
```

##		cities	population	area_km2
## 1		Istanbul	15100000	2576
## 2		Moscow	12500000	2561
## 3		London	9000000	1572
## 4	Saint Petersburg		5400000	1439
## 5		Berlin	3800000	891
## 6		Madrid	3200000	604
## 7		Kyiv	3000000	839
## 8		Rome	2800000	1285
## 9		Bucharest	2200000	228
## 10		Paris	2100000	105

Tibbles

Tibbles are

a **modern reimaging of the data frame**. Tibbles are designed to be (as much as possible) **drop-in replacements** for data frames.

(Wickham, [Advanced R](#))

Have a look at [this book chapter](#) for a full list of the differences between data frames and tibbles and the advantages of using tibbles.

- Tibbles have the same basic properties as data frames (named list of vectors)
- Everything that you can do with data frames, you can do with tibbles



Tibbles

Tibbles are available from the `tibble` package.

Before we use tibbles, we need to install the package once using the function `install.packages`:

```
# This has to be done only once (in the console, not in the script)  
install.packages("tibble")
```

Then, we need to load and attach the package to our current R session using `library`:

```
# This has to be done every time R restarts  
# Put it at the top of a script  
library(tibble)
```



Tibbles

Create a tibble using the `tibble()` function:

```
library(tibble)

tibble(
  cities = cities,
  population = population,
  area_km2 = area_km2
)
```

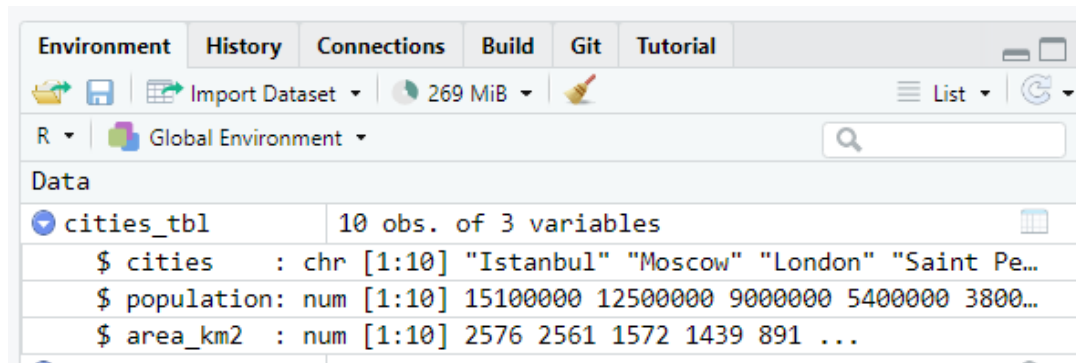
```
## # A tibble: 10 x 3
##   cities      population area_km2
##   <chr>          <dbl>    <dbl>
## 1 Istanbul    15100000    2576
## 2 Moscow      12500000    2561
## 3 London       9000000    1572
## 4 Saint Petersburg 5400000    1439
## 5 Berlin       3800000     891
## 6 Madrid       3200000     604
## 7 Kyiv         3000000     839
## 8 Rome         2800000    1285
## 9 Bucharest    2200000     228
## 10 Paris        2100000     105
```

Exploring tibbles

Look at the structure of an object using `str()`:

```
str(cities_tbl)
## tibble [10 x 3] (S3: tbl_df/tbl/data.frame)
##  $ cities      : chr [1:10] "Istanbul" "Moscow" "London" "Saint Petersburg" ...
##  $ population: num [1:10] 15100000 12500000 9000000 5400000 3800000 3200000 3000000 2800000
##                2200000 2100000
##  $ area_km2   : num [1:10] 2576 2561 1572 1439 891 ...
```

Or do it in the environment pane of RStudio, by clicking the little blue arrow next to the object:



- This works for every R object and is very useful if code doesn't work and you don't know why

Exploring tibbles

How many rows?

```
nrow(cities_tbl)
## [1] 10
```

How many columns?

```
ncol(cities_tbl)
## [1] 3
```

What are the column headers?

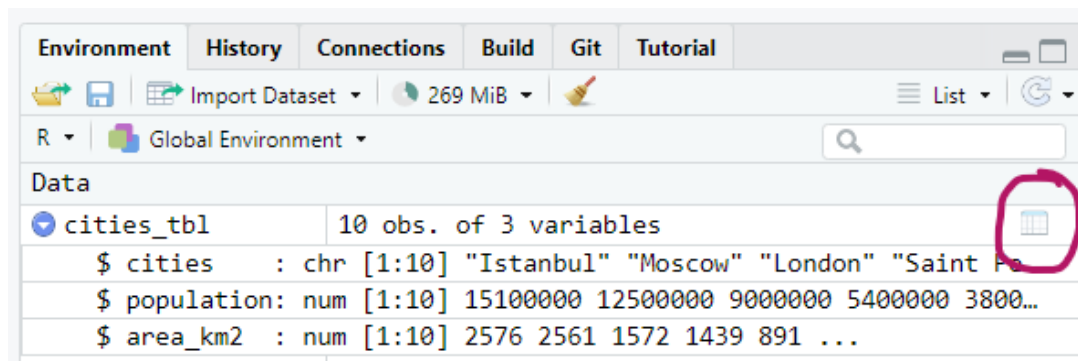
```
names(cities_tbl)
## [1] "cities"      "population" "area_km2"
```

Exploring tibbles

Look at the entire table in a separate window with `view()`:

```
tibble::view(cities_tbl)
```

Or click on the little table sign in the Environment pane:



Exploring tibbles

Get a quick summary of all columns:

```
summary(cities_tbl)
##      cities      population      area_km2
## Length:10      Min.       : 2100000      Min.       : 105.0
## Class :character 1st Qu.: 2850000      1st Qu.: 662.8
## Mode  :character Median   : 3500000      Median   :1088.0
##              Mean    : 5910000      Mean    :1210.0
##              3rd Qu.: 8100000      3rd Qu.:1538.8
##              Max.    :15100000      Max.    :2576.0
```

Indexing tibbles

Indexing tibbles works similar to indexing vectors but with two dimensions instead of 1:

```
tibble [ row_index, col_index or col_name ]
```

- Missing row_index or col_index means *all rows* or *all columns* respectively.
- Indexing a tibble using `[]` always returns another tibble.

Indexing tibbles

```
# First row and first column
cities_tbl[1, 1]
## # A tibble: 1 x 1
##   cities
##   <chr>
## 1 Istanbul
```

This is the same as

```
cities_tbl[1, "cities"]
```

Indexing tibbles: rows

```
# rows 1 & 5, all columns:  
cities_tbl[c(1, 5), ]  
## # A tibble: 2 x 3  
##   cities      population area_km2  
##   <chr>         <dbl>    <dbl>  
## 1 Istanbul    15100000    2576  
## 2 Berlin      3800000     891
```

Indexing tibbles: columns

```
# All rows, first 2 columns
cities_tbl[,1:2] # same as cities_tbl[, c(1, 2)]
# same as
cities_tbl[,c("cities", "population")]
```

```
## # A tibble: 10 x 2
##   cities      population
##   <chr>         <dbl>
## 1 Istanbul    15100000
## 2 Moscow      12500000
## 3 London       9000000
## # ... with 7 more rows
```

Indexing tibbles: columns

Indexing columns by name is usually preferred to indexing by position

```
cities_tbl[ ,1:2] # okay  
cities_tbl[ ,c("cities", "population")] # better
```

Why?

- Code is much easier to read
- Code is more robust against
 - changes in column order
 - mistakes in the code (e.g. typos)

```
cities_tbl[ ,c(1,3)] # 3 instead of 2 -> wrong but no error  
cities_tbl[ ,c("cities", "popluation")] # typo -> wrong and error
```

💡 General rule: Good code produces errors when something unintended or wrong happens

Tibbles: Select columns with \$

Select an entire column from a tibble using `$` (this returns a vector instead of a tibble):

```
cities_tbl$cities
## [1] "Istanbul"      "Moscow"         "London"         "Saint Petersburg"
## [5] "Berlin"        "Madrid"         "Kyiv"           "Rome"
## [9] "Bucharest"     "Paris"
```

Adding new columns

New columns can be added as vectors using the `$` operator. The vectors need to have the same length as the tibble has rows.

```
# add a country column
cities_tbl$country <- c(
  "Turkey", "Russia", "UK", "Russia", "Germany", "Spain",
  "Ukraine", "Italy", "Romania", "France"
)
```

```
## # A tibble: 10 x 4
##   cities      population area_km2 country
##   <chr>      <dbl>    <dbl> <chr>
## 1 Istanbul    15100000    2576 Turkey
## 2 Moscow      12500000    2561 Russia
## 3 London       9000000    1572 UK
## 4 Saint Petersburg 5400000    1439 Russia
## 5 Berlin       3800000     891 Germany
## 6 Madrid       3200000     604 Spain
## 7 Kyiv         3000000     839 Ukraine
## 8 Rome         2800000    1285 Italy
## 9 Bucharest    2200000     228 Romania
## 10 Paris       2100000     105 France
```

Summary I

data frames and tibbles

- can be used to represent tables in R
- are pretty similar, however tibbles are slightly convenient and modern
- are **named lists of vectors of the same length**
 - every column is a vector
 - columns have a header which is the name of the vector in the list
 - within one column, values are of same data type
 - every column has the same length

tibbles

- to use tibbles, install the package once with `install.packages("tibble")`
- put `library(tibble)` at the beginning of your script to load package

Summary II

Creating tibbles and data frames

```
# data frame
data.frame(
  a = 1:3,
  b = c("a", "b", "c"),
  c = c(TRUE, FALSE, FALSE)
)
# tibble
tibble(
  a = 1:3,
  b = c("a", "b", "c"),
  c = c(TRUE, FALSE, FALSE)
)
# convert data frame to tibble
as_tibble(df)
```

Summary III

Looking at tibble structure

```
# structure of tibble and data types of columns  
str(tbl)  
# number of rows  
nrow(tbl)  
# number of columns  
ncol(tbl)  
# column headers  
names(tbl)  
# look at the data in a new window  
tibble::view(tbl)  
# summary of values from each column  
summary(tbl)
```

Summary IV

Indexing tibbles and selecting columns

Return result as tibble:

```
# rows and columns by position
tbl[1:3, c(1, 3)]
tbl[1:3, ] # all columns
tbl[, 3] # column 3, all rows
tbl[3] # same as above

# columns by name
tbl[, c("colA", "colB")]
tbl[c("colA", "colB")]
```

Return result as vector:

```
tbl$colA # select colA
```

Now you

Task 3: Tibbles (20 min)

Find the task description [here](#)