

# Statistical tests


## Introduction to R - Day 3

Instructor: **Selina Baldauf**

Freie Universität Berlin - Theoretical Ecology

2021-08-01 (updated: 2023-02-26)

# Overview of tests



Please note the discussion going on about p-values as a basis for binary decisions. See e.g. [here](#) or [here](#) as a starting point

# Tests for normal distribution

# Test for normal distribution

There are **various tests** and the outcome might differ!

## Shapiro-Wilk-Test

- How much does variance of observed data differ from normal distribution
- Specific test only for normal distribution
- High power, also for few data points

## Visual tests: QQ-Plot


- Quantiles of observed data plotted against quantiles of normal distribution
- Scientist has to decide if normal or not

# The data

Create a tibble with two variables

- `normal`: 200 normally distributed values with mean 50 and standard deviation 5
- `non_normal`: 200 uniformly distributed values between 45 and 55

```
set.seed(123)
mydata <- tibble(
  normal = rnorm(
    n = 200,
    mean = 50,
    sd = 5
  ),
  non_normal = runif(
    n = 200,
    min = 45,
    max = 55
  )
)
```



# Shapiro-Wilk-Test

$H_0$ : Data does not differ from a normal distribution

```
shapiro.test(mydata$normal)
##
##      Shapiro-Wilk normality test
##
## data:  mydata$normal
## W = 0.99076, p-value = 0.2298
```

- W: test statistic
- p-value: probability to observe the data if  $H_0$  was true

The data does not deviate significantly from a normal distribution (Shapiro-Wilk-Test,  $W = 0.991$ ,  $p = 0.23$ ).


```
shapiro.test(mydata$non_normal)
##
##      Shapiro-Wilk normality test
##
## data:  mydata$non_normal
## W = 0.95114, p-value = 2.435e-06
```

The data deviates significantly from a normal distribution (Shapiro-Wilk-Test,  $W = 0.95$ ,  $p < 0.001$ ).


# Visual test with QQ-Plot

Points should match the straight line. Small deviations are okay.

```
# ggplot(mydata, aes(sample = normal)) +  
#   stat_qq() + stat_qq_line()  
ggpubr::ggqqplot(mydata$normal)
```



```
# ggplot(mydata, aes(sample = non_normal)) +  
#   stat_qq() + stat_qq_line()  
ggpubr::ggqqplot(mydata$non_normal)
```



# Tests for equal variance




# The data

Counts of insects in agricultural units treated with different insecticides.

Compare treatments A, B and C:

- Create subsets before: count variable for each treatment as a vector

```
TreatA <- filter(InsectSprays,  
                  spray == "A")$count  
TreatB <- filter(InsectSprays,  
                  spray == "B")$count  
TreatC <- filter(InsectSprays,  
                  spray == "C")$count
```



# Test for equal variance

First, test for normal distribution!

## F-Test

- **Normal distribution** of groups
- Calculates ratio of variances (if equal, ratio = 1)
- p: How likely is ratio if variances were equal?

## Levene test

- **Non-normal distribution** of groups
- Compare difference between data sets with difference within data sets

# Test for equal variances

If we want to compare variances between treatments A, B and C, we first test for normal distribution

```
shapiro.test(TreatA)
##
##      Shapiro-Wilk normality test
##
## data:  TreatA
## W = 0.95757, p-value = 0.7487
shapiro.test(TreatB)
##
##      Shapiro-Wilk normality test
##
## data:  TreatB
## W = 0.95031, p-value = 0.6415
shapiro.test(TreatC)
##
##      Shapiro-Wilk normality test
##
## data:  TreatC
## W = 0.92128, p-value = 0.2967
```

Result: All 3 treatments are normally distributed.

# F-Test

$H_0$ : Variances do not differ between groups

```
var.test(TreatA, TreatB)
##
##      F test to compare two variances
##
## data:  TreatA and TreatB
## F = 1.2209, num df = 11, denom df = 11, p-value = 0.7464
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
##  0.3514784 4.2411442
## sample estimates:
## ratio of variances
##           1.22093
```

- F: test statistics, ratio of variances (if  $F = 1$ , variances are equal)
- df: degrees of freedom of both groups
- p-value: how likely is it to observe the data if  $H_0$  was true?

Result: The variances of sprays A and B do not differ significantly (F-Test,  $F_{11,11} = 1.22$ ,  $p = 0.75$ )

# F-Test

$H_0$ : Variances do not differ between groups

```
var.test(TreatA, TreatC)
##
##      F test to compare two variances
##
## data:  TreatA and TreatC
## F = 7.4242, num df = 11, denom df = 11, p-value = 0.002435
## alternative hypothesis: true ratio of variances is not equal to 1
## 95 percent confidence interval:
##   2.137273 25.789584
## sample estimates:
## ratio of variances
##           7.424242
```

Result: The variances of sprays A and C differ significantly (F-Test,  $F_{11,11} = 7.42$ ,  $p = 0.002$ )

# Test for equal means

# Test for equal means

## t-test


- Normal distribution AND equal variance
- Compares if mean values are within range of standard error of each other
- p: how likely is the difference if the means were equal

## Welch-Test

- Normal distribution but unequal variance
- Corrected t-test

## Wilcoxon rank sum test

- Non-normal distribution and unequal variance
- Compares rank sums of the data
- Non-parametric



# t-test

$H_0$ : The samples do not differ in their mean

Treatment A and B: **normally distributed** and **equal variance**

```
t.test(TreatA, TreatB, var.equal = TRUE)
##
##      Two Sample t-test
##
## data:  TreatA and TreatB
## t = -0.45352, df = 22, p-value = 0.6546
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  -4.643994  2.977327
## sample estimates:
## mean of x mean of y
##  14.50000  15.33333
```

- t: test statistics (t = 0 means equal means)
- df: degrees of freedom of t-statistics
- p-value: how likely is it to observe the data if  $H_0$  was true?

**Result:** The means of spray A and B do not differ significantly (t = -0.45, df = 22, p = 0.66)



# Welch-Test

$H_0$ : The samples do not differ in their mean

Treatment A and C: **normally distributed** and **non-equal variance**

```
t.test(TreatA, TreatC, var.equal = FALSE)
##
##      Welch Two Sample t-test
##
## data:  TreatA and TreatC
## t = 7.5798, df = 13.91, p-value = 2.655e-06
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##    7.885546 14.114454
## sample estimates:
## mean of x mean of y
##    14.5      3.5
```

**Result:** The means of spray A and C do differ significantly ( $t = 7.58$ ,  $df = 13.9$ ,  $p < 0.001$ )

# Wilcoxon-rank-sum Test

$H_0$ : The samples do not differ in their mean

We don't need the Wilcoxon test to compare treatment A and B, but for the sake of an example:

```
wilcox.test(TreatA, TreatB)
##
##      Wilcoxon rank sum test with continuity correction
##
## data:  TreatA and TreatB
## W = 62, p-value = 0.5812
## alternative hypothesis: true location shift is not equal to 0
```

**Result:** The means of spray A and B do not differ significantly ( $W = 62$ ,  $p = 0.58$ )

# Paired values

Are there pairs of data points?

**Example:** samples of invertebrates across various rivers before and after sewage plants.

- For each plant, there is a pair of data points (before and after the plant)
- Question: Is the change (before-after) significant

Use `paired = TRUE` in the test.

```
t.test(TreatA, TreatB, var.equal = TRUE, paired = TRUE)
t.test(TreatA, TreatB, var.equal = FALSE, paired = TRUE)
wilcox.test(TreatA, TreatB, paired = TRUE)
```

Careful: your treatment vector both have to have the same order

# Plot test results with `ggsignif`


The `ggsignif` package offers a `geom_signif()` layer that can be added to a ggplot to annotate significance levels

```
# install.packages("ggsignif")  
library(ggsignif)
```

# Plot test results with `geom_signif()`

```
ggplot(InsectSprays,
      aes(x = spray, y = count)) +
  geom_boxplot(notch = TRUE) +
  geom_signif(
    comparisons = list(
      c("A", "B"),
      c("B", "C"),
      c("A", "C")
    ),
    map_signif_level = TRUE,
    y_position = c(23, 24, 25)
  )
```


- By default, a Wilcoxon test is performed



# Plot test results with `geom_signif()`

```
ggplot(InsectSprays,
      aes(x = spray, y = count)) +
  geom_boxplot(notch = TRUE) +
  geom_signif(
    comparisons = list(
      c("A", "B"),
      c("B", "C"),
      c("A", "C")
    ),
    test = "t.test",
    test.args = list(
      var.equal = TRUE
    ),
    map_signif_level = TRUE,
    y_position = c(23, 24, 25)
  )
```

- Specify a specific test with the `test` argument
- Additional arguments can be passed as a list with the `test.args` argument
- Look at `?geom_signif` for more options




# Plot mean $\pm$ se using `stat_summary`

Another way to plot the results is to plot mean and standard error of the mean:

```
ggplot(  
  InsectSprays,  
  aes(x = spray, y = count)  
) +  
  stat_summary()
```

- By default `stat_summary` adds mean and standard error of the mean as pointrange




# Plot mean $\pm$ se using `stat_summary`

Another way to plot the results is to plot mean and standard error of the mean:

```
ggplot(InsectSprays, aes(x = spray, y = count))  
+  
  stat_summary(  
    fun.data = mean_se,  
    geom = "errorbar"  
  ) +  
  stat_summary(  
    fun.y = mean,  
    geom = "point",  
    color = "#28a87d",  
    size = 4  
  )
```

- Inside `stat_summary`, you can define function used to calculate the summary
  - `fun.data` for errorbars
  - `fun.y` for point values (e.g. mean)
- Define the geom that represents the summary






# Plot mean $\pm$ se using `stat_summary`

Another way to plot the results is to plot mean and standard error of the mean:


```
ggplot(InsectSprays, aes(x = spray, y = count))  
+  
  stat_summary(  
    fun.data = mean_se,  
    geom = "errorbar",  
    width = 0.3  
  ) +  
  stat_summary(  
    fun.y = mean,  
    geom = "bar",  
    size = 4  
  )
```



# Plot mean $\pm$ se using `stat_summary`

Just like before, you can also add a `geom_signif` to a barplot:

```
ggplot(InsectSprays, aes(x = spray, y = count))
+
  stat_summary(
    fun.data = mean_se,
    geom = "errorbar",
    width = 0.3
  ) +
  stat_summary(
    fun.y = mean,
    geom = "bar"
  ) +
  ggsignif::geom_signif(
    comparisons = list(
      c("A", "B"),
      c("B", "C"),
      c("A", "C")
    ),
    test = "t.test",
    map_signif_level = TRUE,
    y_position = c(17, 18, 19)
  )
```



# Now you

Task 1: Statistical tests (45 min)

Find the task description [here](#)